

# Towards Generating Coherent and Diverse Text with Deep Generative Models

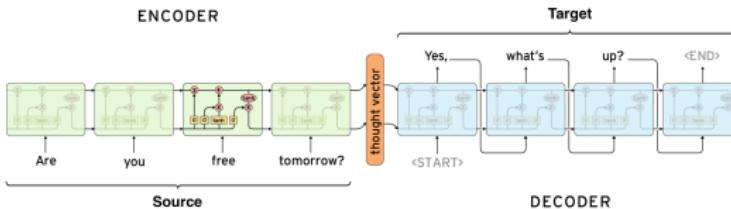
**Presenter:** Yizhe Zhang

Duke University

October 25, 2017

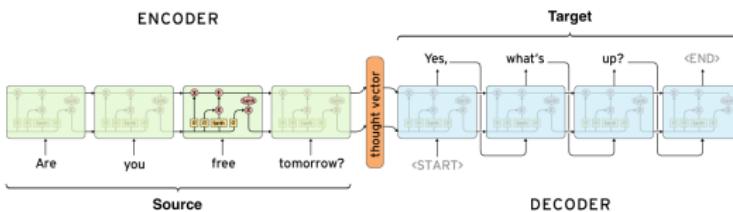
# Text generation

- **Human:** ideas → utterance
- **Machine:** latent representation → understandable sentences.
- Language modeling, translation, summarization, conversation models, image/video captioning...
- MLE-based modeling with RNN (esp. LSTM) achieves success.
- Variants: Hierarchical, attentive, etc...



# Challenges in text generation with standard model

- MLE-based(e.g. seq2seq) model with RNN has several issues:
  - *Teacher forcing*: the discrepancy between training and testing → inconsistency
  - *Exposure bias*: error aggregates with length. → generation quality decreases.
  - *MLE-based*: lose distributional information. → underestimate diversity and fail to capture multiple modes.



## Related works

- Efforts had been made in alleviating above issues
- For *teacher forcing* and *exposure bias* issues
- Scheduled Sampling [Bengio, et al., NIPS 2015]
- Minimum Risk Training [Shen et al., ACL 2016]
- Professor Forcing [Alex et. al, NIPS 2016]
- RL methods:
  - MIXER [Ranzato et al., ICLR 2016]
  - Actor-Critic [Bahdanau et al., ICLR 2017]
  - SeqGAN [Yu et al., AAAI 2017]

# Why Deep generative models for text ?

- Deep generative models (DGMs): generative adversarial network (GAN), variational autoencoder (VAE)
- Probabilistic framework account for uncertainty and diversity, natural recipe for generation.
- Ideally, can be leveraged to alleviate *teacher forcing* and *exposure bias* issues.
- Previous works
  - RNN-VAE[Bowman et al., Arxiv 2015]
  - SeqGAN [Yu et al., AAAI 2017]
  - Actor-Critic [Bahdanau et al., ICLR 2017]

Generating realistic text from given corpus

# Outline

## 1 Overview

## 2 Generating text via adversarial training

- Generating realistic text from given corpus
- Generating coherent and diverse conversation

## 3 Deconvolutional text modeling

- Deconvolutional text autoencoder
- Deconvolutional text VAE

## 4 Summary







Generating realistic text from given corpus

# Feature moment matching

Overview

- The adversarial game is the following:
- $D(\cdot)$  attempts to select informative sentence features.
- $G(\cdot)$  aims to match these features.
- Features are selected according to **syn/real discrimination ability, latent code reconstruction and moment matching precision**.

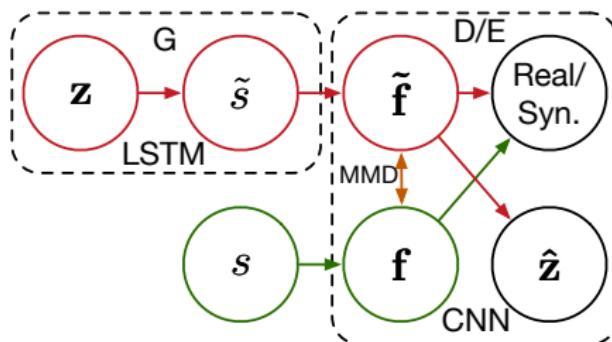


Figure: Model scheme of TextGAN.

Generating realistic text from given corpus

# Feature moment matching

Generator

- Optimization schemes:

$$\mathcal{L}_G = \mathcal{L}_{MMD^2}$$

$$\mathcal{L}_D = \mathcal{L}_{GAN} + \lambda_r \mathcal{L}_{recon} - \lambda_m \mathcal{L}_{MMD^2}$$

- For G, consider a moment matching loss over *feature vector* using *maximum mean discrepancy* (MMD).

$$\begin{aligned}\mathcal{L}_{MMD^2} &= \|\mathbb{E}_{x \sim \mathcal{X}} \phi(x) - \mathbb{E}_{y \sim \mathcal{Y}} \phi(y)\|_{\mathcal{H}}^2 \\ &= \mathbb{E}_{x \sim \mathcal{X}} \mathbb{E}_{x' \sim \mathcal{X}} [k(x, x')] \\ &\quad + \mathbb{E}_{y \sim \mathcal{Y}} \mathbb{E}_{y' \sim \mathcal{Y}} [k(y, y')] - 2 \mathbb{E}_{x \sim \mathcal{X}} \mathbb{E}_{y \sim \mathcal{Y}} [k(x, y)].\end{aligned}\tag{4}$$

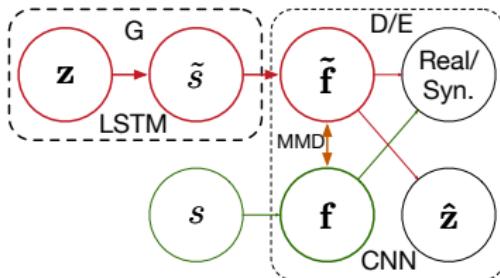
- With a Gaussian kernel , minimizing the MMD objective  $\Leftrightarrow$  minimizing *all order of moments* of two empirical distributions.

Generating realistic text from given corpus

# Feature moment matching

Generator

- Vanilla GAN: D independently judge each syn/real data.
- The **MMD loss** for G: match distributions, enforce diversity.
- The gradient signal back-propagated from feature layer is more direct.



Generating realistic text from given corpus

# Feature moment matching

Discriminator

- Optimization schemes:

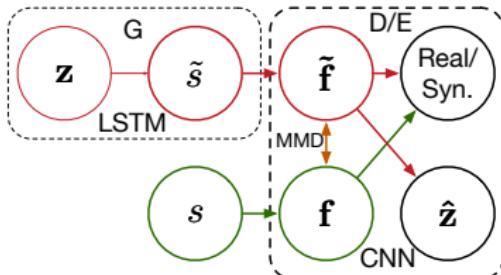
$$\mathcal{L}_G = \mathcal{L}_{MMD^2}$$

$$\mathcal{L}_D = \mathcal{L}_{GAN} + \lambda_r \mathcal{L}_{recon} - \lambda_m \mathcal{L}_{MMD^2}$$

$$\mathcal{L}_{GAN} = -\mathbb{E}_{s \sim \mathcal{S}} \log D(s) - \mathbb{E}_{z \sim p_z} \log [1 - D(G(z))]$$

$$\mathcal{L}_{recon} = \|\hat{z} - z\|^2,$$

- The **reconstruction loss** in D : select the most *representative* (information-preserving) features.
- The **MMD loss** in D: select the most *challenging* features.



Generating realistic text from given corpus

# Framework components

## LSTM generator

- We specify an LSTM generator to translate a *latent code vector*,  $z$ , into a *synthetic sentence*  $\tilde{s}$ .
- All other words in the sentence are sequentially generated using the RNN, *based on previously generated words*, until the end-sentence symbol is generated.

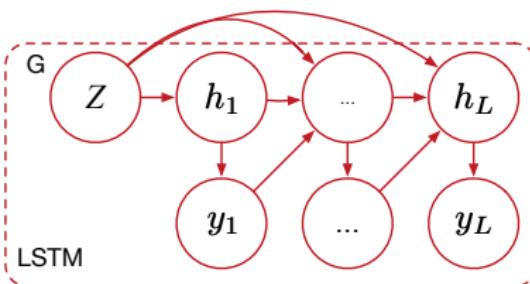


Figure: LSTM generator

Generating realistic text from given corpus

# Framework components

## CNN discriminator

- CNNs weight each word equally and are empirically better at abstracting features particularly with long sentences.
- A sentence is represented as a matrix  $\mathbf{X} \in \mathbb{R}^{k \times T}$ , followed by a convolution operation.
- A max-over-time pooling operation is then applied.

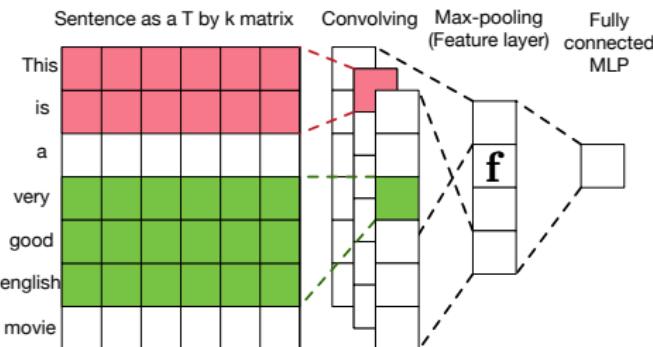


Figure: CNN discriminator

Generating realistic text from given corpus

## Empirical evaluation

- **Dataset:** 0.5M Arxiv sentences + 0.5M BookCorpus sentences .
- **Evaluation:** Kernel density estimation (KDE) .
- **Evaluation:** Corpus-level BLEU score .
- **Comparison:** baseline auto-encoder, variational auto-encoder and seqGAN [Yu et al. 2016]

Generating realistic text from given corpus

# Experimental Result

Generated text

- Produce novel phrases. (b)
- The synthetic sentences seem syntactically reasonable.
- The semantic meaning is less preserved. (e)

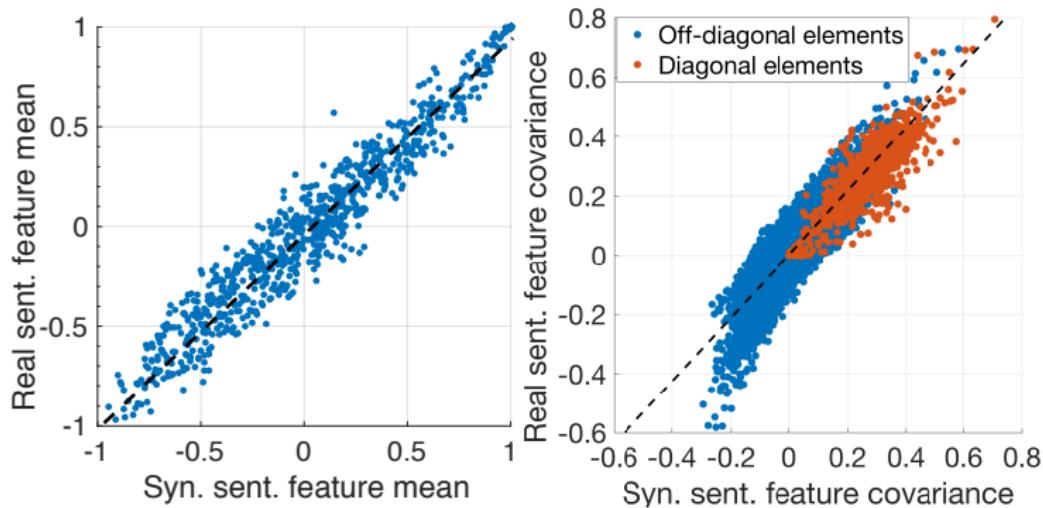
Table: Sentences generated by textGAN.

- 
- a we show the joint likelihood estimator ( in a large number of estimating variables embedded on the subspace learning ) .
  - b this problem achieves less interesting choices of convergence guarantees on turing machine learning .
  - c in hidden markov relational spaces , the random walk feature decomposition is unique generalized parametric mappings.
  - d i see those primitives specifying a deterministic probabilistic machine learning algorithm .
  - e i wanted in alone in a gene expression dataset which do n't form phantom action values .
  - f as opposite to a set of fuzzy modelling algorithm , pruning is performed using a template representing network structures .
-

Generating realistic text from given corpus

# Experimental Result

## Moment Matching



**Figure:** Moment matching comparison. Left: expectations of latent features from real vs. synthetic data. Right: elements of covariance matrix for real and synthetic data, respectively.

Generating realistic text from given corpus

# Experimental Result

## Sentence transition

**Table:** Intermediate sentences by linear transition.

	textGAN	AE
<b>A</b>	our methods apply novel approaches to solve modeling tasks .	
- our methods apply novel approaches to solve modeling .	our methods apply to train UNK models involving complex .	
- our methods apply two different approaches to solve computing .	our methods solve use to train ) .	
- our methods achieves some different approaches to solve computing .	our approach show UNK to models exist .	
- our methods achieves the best expert structure detection .	that supervised algorithms show to UNK speed .	
- the methods have been different related tasks .	that address algorithms to handle ) .	
- the guy is the minimum of UNK .	that address versions to be used in .	
- the guy is n't easy tonight .	i believe the means of this attempt to cope .	
- i believe the guy is n't smart okay?	i believe it 's we be used to get .	
- i believe the guy is n't smart .	i believe it i 'm a way to belong .	
<b>B</b>	i believe i 'm going to get out .	

Generating realistic text from given corpus

# Experimental Result

## Quantitative evaluation

- Higher BLEU, lower KDE is better.

Table: Quantitative results using BLEU-2,3,4 and KDE.

	BLEU-4	BLEU-3	BLEU-2	KDE(nats)
AE	0.01±0.01	0.11±0.02	0.39±0.02	2727±42
VAE	0.12±0.06	0.40±0.06	0.61±0.07	2025±25
seqGAN	0.04±0.04	0.30±0.08	0.67±0.04	2019±53
textGAN(MM)	0.09±0.04	0.42±0.04	0.77±0.03	1823±50
textGAN(CM)	0.12±0.03	0.49±0.06	0.84±0.02	1686±41
textGAN(MMD)	<b>0.13±0.05</b>	0.49±0.06	0.83±0.04	1688±38
textGAN(MMD-L)	0.11±0.05	<b>0.52±0.07</b>	<b>0.85±0.04</b>	<b>1684±44</b>

Generating realistic text from given corpus

# Summary

- A framework for text generation using adversarial training
- Techniques to alleviate practical issues when training GAN on text domain.
- Future works:
  - Variance reduction for gradient estimation.
  - Better approach for reducing exposure bias
- code available at:  
[https://github.com/dreasysnail/textGAN\\_public](https://github.com/dreasysnail/textGAN_public)

Generating coherent and diverse conversation

# Outline

## 1 Overview

## 2 Generating text via adversarial training

- Generating realistic text from given corpus
- Generating coherent and diverse conversation

## 3 Deconvolutional text modeling

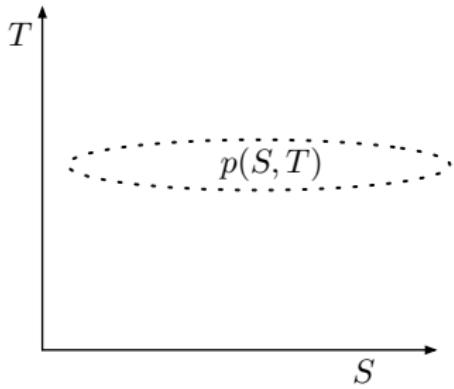
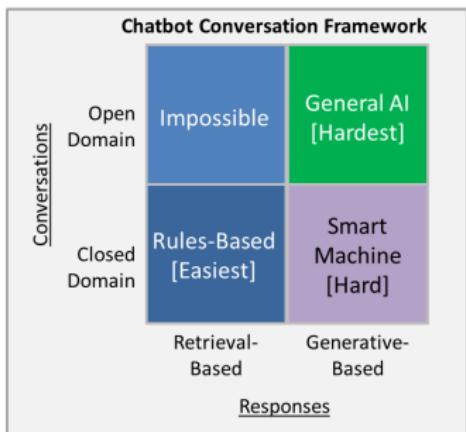
- Deconvolutional text autoencoder
- Deconvolutional text VAE

## 4 Summary

Generating coherent and diverse conversation

# Open-domain generative neural conversation model

- Single turn conversation.  $S$ : source prompt,  $T$ : target response
- Seq2seq MLE:  $\hat{T} = \operatorname{argmax}_T \log p(T|S)$
- **Issue:** bland and not diverse responses: I do n't know



Generating coherent and diverse conversation

## Previous methods

- Maximum mutual information (MMI) [Li, et al., NAACL 2016]:

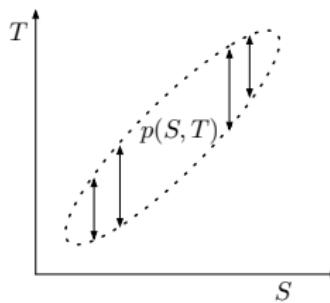
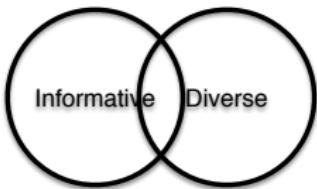
$$\hat{T} = \operatorname{argmax}_T \log p(T|S) - \lambda \log p(T)$$

- **Issue:** 1) hyperparameter makes it less principled. 2) training/testing discrepancy

Generating coherent and diverse conversation

# Diversity and blandness

- *Diversity*: variance of  $p(T|S)$  with fixed  $S$
- *Blandness*: variation of  $p(T|S)$  w.r.t.  $S$
- Diverse but bland:  
I do n't know; Not sure; We'll see; I agree
- Informative but monotonous: deterministic response



Generating coherent and diverse conversation

## Our aim

- Improve *diversity* and reduce *blandness*
- *diversity*: use adversarial training to fit conditional distribution  $P(T|S)$ .
- *blandness*: use variational information maximization (VIM) to generate informative response.

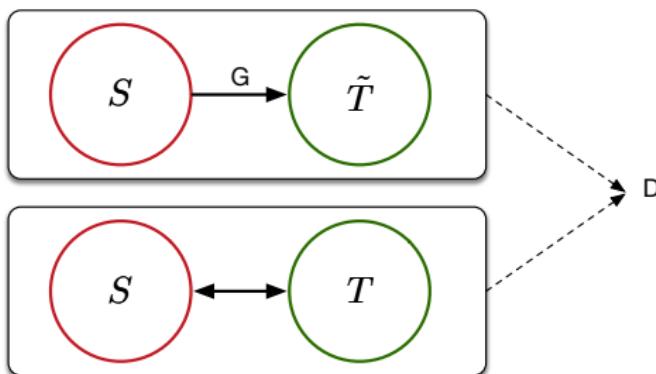
Generating coherent and diverse conversation

# Conditional GAN

## Objective

- Using conditional GAN:

- ① Rather than model  $\text{argmax}_T p(T|S)$ , consider  $p(T|S)$ .
- ② Sequence-level loss, alleviating exposure bias.



Generating coherent and diverse conversation

# Conditional GAN

## Generator

- $(S, z) \mapsto \hat{T}$ , CNN-LSTM seq2seq model.
- *Gumbel-softmax* instead of policy gradient.
- Truncated back-propagation.

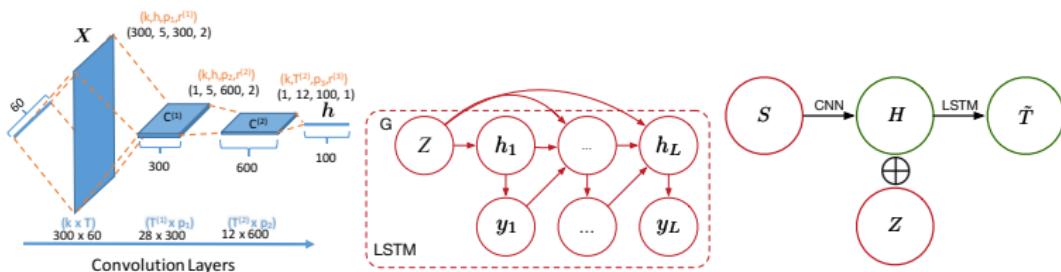


Figure: CNN-LSTM generator

Generating coherent and diverse conversation

# Conditional GAN

## Discriminator

*Sent.feature* :  $E_s \triangleq E_{w_s}(S), E_t \triangleq E_{w_t}(T)$

*Discr.output* :  $D((T, S)) = D_{cos}(E_{w_s}(S), E_{w_t}(T))$

- Inspired by DSSM [Huang et al., CIKM 2013]

$$\mathcal{L}_G = \mathbb{E}_{(T, \tilde{T}, S) \sim p, p^e} f'(D((T, S)) - D((\tilde{T}, S))) , \quad \mathcal{L}_D = -\mathcal{L}_G$$

Where  $f' : [-2, 2] \mapsto \mathbb{R}$

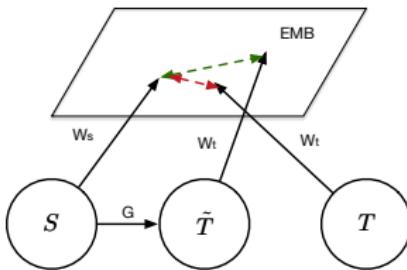
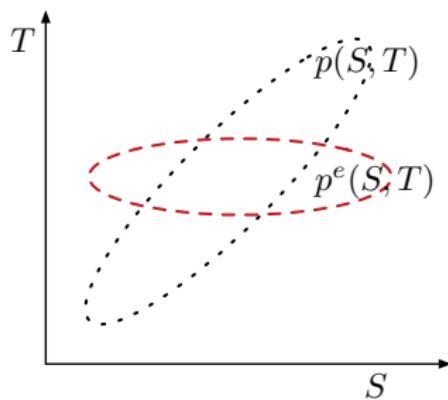


Figure: Discriminator

## Generating coherent and diverse conversation

## Blandness problem

- *Oracle joint distr.*:  $p(T, S)$
- *Encoder joint distr.*:  $p_\theta^e(T, S) = p_\theta^e(T|S)p(S)$
- Our aim: learn  $\theta$ , such that  $p_\theta^e(T, S)$  approximate  $p(T, S)$
- **Problem:** learned  $p_\theta^e(T|S)$  tends to be similar with different  $S$ .



## Generating coherent and diverse conversation

## Maximize mutual information

- **Solution:** explicitly maximize mutual information  $I_{p^e}(S, T)$  over  $p^e$ .
- $I_{p^e}(S, T)$  is hard to directly optimize
- Use *variational information maximization* (VIM):
- A lower bound of  $I_{p^e}(S, T)$

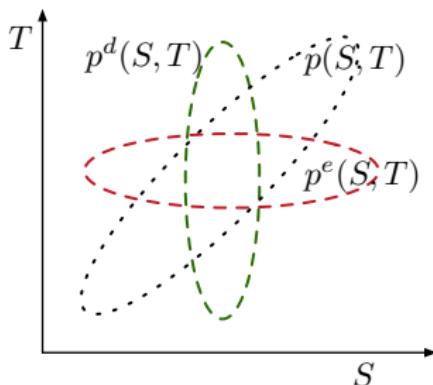
$$\begin{aligned}
 I_{p^e}(S, T) &= \mathbb{E}_{p^e(S, T)} \log \frac{p^e(S, T)}{p(S)p(T)} \\
 &= -\mathbb{E}_{p(S)} \log p(S) + \mathbb{E}_{p(T)} \mathbb{E}_{p^e(S|T)} \log p^e(S|T) \\
 &\geq \mathbb{E}_{p(S)} \mathbb{E}_{p_\theta(T|S)} \log q_\phi(S|T) + \text{Constant}
 \end{aligned} \tag{5}$$

- Different from Li et al. 2015, where forward and backward models are separable.

Generating coherent and diverse conversation

# Learn a better joint distribution

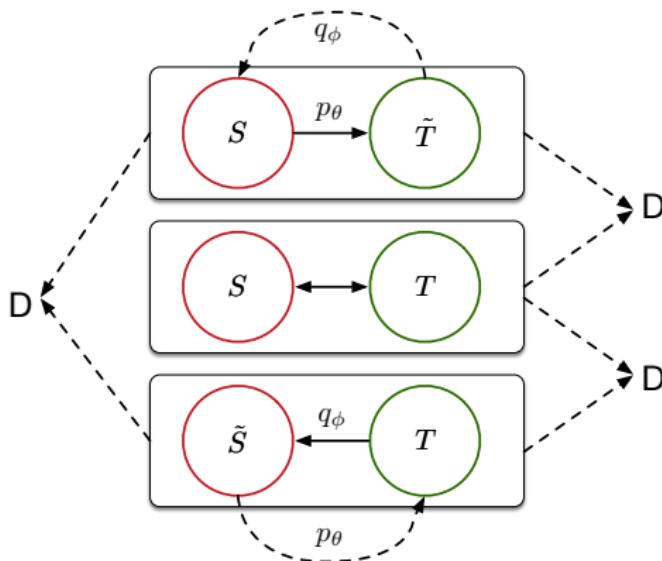
- $q_\phi(S|T)$  may also tend to be bland in generating source.



- Forward and backward model are *symmetric* and *collaborative*.
- What about make the objective symmetric?
- *Decoder joint distr.*:  $p_\theta^d(T, S) = q_\phi^d(S|T)p(T)$

Generating coherent and diverse conversation

# Learn a better joint distribution (Cont'd)



- Augmented paired data for discriminator training
- Forward and backward model works in a synergistic manner

Generating coherent and diverse conversation

# Final Objective

- Joint train a supervised loss, similar to Xia et. al, ICML 2017
- Final objective:

$$\begin{aligned}
 & \operatorname{argmin}_{\psi} \operatorname{argmax}_{\theta, \phi} \mathcal{L} \\
 &= -\mathbb{E}_{(T, \tilde{T}, S) \sim p, p_{\theta}^e} f'(D_{\psi}((T, S)) - D_{\psi}((\tilde{T}, S))) \\
 &\quad - \mathbb{E}_{(T, \tilde{S}, S) \sim p, p_{\phi}^d} f'(D_{\psi}((T, S)) - D_{\psi}((T, \tilde{S}))) \\
 &\quad + \lambda \mathbb{E}_{p(S)} \mathbb{E}_{p_{\theta}(T|S)} \log q_{\phi}(S|T) \\
 &\quad + \lambda \mathbb{E}_{p(T)} \mathbb{E}_{q_{\phi}(S|T)} \log p_{\theta}(T|S)
 \end{aligned} \tag{6}$$

where  $\lambda$  is a balance parameter.

Generating coherent and diverse conversation

# Experiments

- Reddit 2M dataset, src-tgt pairs.
- Maximum length 53.
- Use BLEU-4, entropy-4, distinct-1,2 for relevancy and diversity evaluation.
- Distinct: (# of unique token)/(# of total token)
- Entropy:

$$E = -\frac{1}{T} \sum_w F(w) [\log F(w) - \log T] \quad (7)$$

$$T = \sum_w F(w) \quad (8)$$

where  $F(w)$  denote the occurrence of token  $w$ .

Generating coherent and diverse conversation

# Experimental Result

## Quantitative evaluation

- Higher BLEU, entropy and distinct value is better.

Table: Quantitative results

	BLEU-4	Dist-1	Dist-2	Ent-4
Seq2seq	1.77	0.005	0.041	7.421
CondGAN	1.85	0.008	0.072	9.679
Dual-GAN+VIM	<b>2.32</b>	<b>0.009</b>	<b>0.093</b>	<b>10.213</b>

Generating coherent and diverse conversation

# Experimental Result

Generated responses

- Random questions

<b>Source:</b>	How do you think of seattle ?
<b>Seq2seq:</b>	i was n't thinking about it .
<b>CondGAN:</b>	i think it 's not that i want to know .
<b>Bidir-GAN</b>	it 's a great place, i 'd say to visit !

- CondGAN

<b>Source:</b>	which country do you want to travel ?
<b>Seq2seq:</b>	american , but i 'm not sure what the country is .
<b>CondGAN:</b>	i 'm not sure i know of any country , but i 'm just saying what the other is for .
<b>Bidir-GAN</b>	i think it 's called " south africa ".

Generating coherent and diverse conversation

# Summary

- Summary

- Adversarial seq2seq training scheme with specially designed sequence-level discrimination loss.
- Simultaneously learning forward and backward models, which enhance each other.

- Future work:

- Theoretical perspective for the adversarial loss.
- Stabilize GAN training.

Deconvolutional text autoencoder

# Outline

## 1 Overview

## 2 Generating text via adversarial training

- Generating realistic text from given corpus
- Generating coherent and diverse conversation

## 3 Deconvolutional text modeling

- Deconvolutional text autoencoder
- Deconvolutional text VAE

## 4 Summary

Deconvolutional text autoencoder

# Fully convolutional text autoencoder

- *RNN as generator*: sequential generation process. → exposure bias, even with GAN...
- Different decoder/generator?
- Recent work: DCNN+RNN [Semeniuta, et al., 2017, arxiv], PixelCNN [Yang, et al., 2017, arxiv] → *still sequential model*
- *Purely convolutional/transpose-convolutional framework* for decoder/generator. [Zhang et al., NIPS 2017]

## Deconvolutional text autoencoder

## Fully convolutional text autoencoder

- Consider a text-autoencoder with CNN as encoder and deconvolutional network (DCNN) as decoder

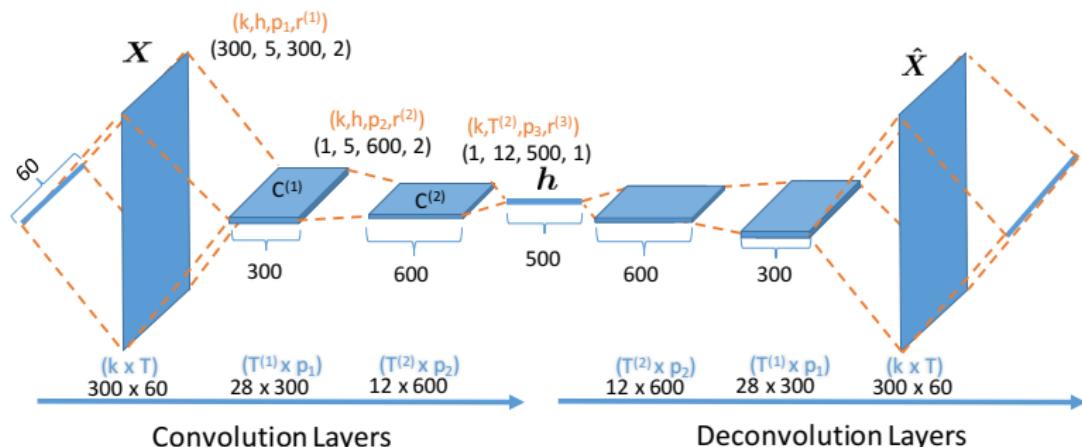


Figure: Framework

## Deconvolutional text autoencoder

## Model details

- *Input* sentence of length  $T$ :  $\mathbf{X} \in \mathbb{R}^{k \times T}$
- *Embedding matrix*:  $\mathbf{W}_e \in \mathbb{R}^{k \times V}$ .
- The *CNN* encodes sentence  $s$  into a vector  $\mathbf{h}$

$$\mathbf{h} = \text{multi-layer-conv}(\mathbf{X}) \quad (9)$$

- The *DCNN* decodes  $\mathbf{h}$  into  $\hat{\mathbf{Y}}$

$$\hat{\mathbf{X}} = \text{multi-layer-deconv}(\mathbf{h}) \quad (10)$$

- Let  $\hat{w}^t$  denote the  $t$ -th word in *generated* sentence  $\hat{s}$ , the probability of  $\hat{w}^t$  to be a specific word  $v$  is

$$p(\hat{w}^t = v) = \frac{\exp[D_{cos}(\hat{\mathbf{x}}^t, \mathbf{W}_e^v)/\tau]}{\sum_{v' \in V} \exp[D_{cos}(\hat{\mathbf{x}}^t, \mathbf{W}_e^{v'})/\tau]} \quad (11)$$

$D_{cos}(\mathbf{x}, \mathbf{y})$  denotes cosine similarity,  $\tau$  is a temperature parameter ( $\tau = 0.01$ ).

# DCNN as a decoder/generator

- Consistent training and testing → *free of teacher forcing*
- Without a pre-specified ordering structure → *free of exposure bias*
- *Fast* and easy to be *parallelized*.
- *Weaker* dependency, less order-preserving, better for higher mutual information source-target pairs.
- scenarios for applying DCNN decoder rather than LSTM decoder:  
*Reconstruction > summarization > translation > conversation*

## Deconvolutional text autoencoder

## Reconstruction

Ground-truth:	on every visit to nyc , the hotel beacon is the place we love to stay . so conveniently located to central park , lincoln center and great local restaurants . the rooms are lovely . beds so comfortable , a great little kitchen and new wizz bang coffee maker . the staff are so accommodating and just love walking across the street to the fairway supermarket with every imaginable goodies to eat .
Hier. LSTM [Li et al, 2015]	every time in new york , lighthouse hotel is our favorite place to stay . very convenient , central park , lincoln center , and great restaurants . the room is wonderful , very comfortable bed , a kitchenette and a large explosion of coffee maker . the staff is so inclusive , just across the street to walk to the supermarket channel love with all kinds of what to eat .
Our CNN- DCNN	on every visit to nyc , the hotel beacon is the place we love to stay . so closely located to central park , lincoln center and great local restaurants . biggest rooms are lovely . beds so comfortable , a great little kitchen and new UNK suggestion coffee maker . the staff turned so accommodating and just love walking across the street to former fairway supermarket with every food taxes to eat .

Table: Reconstructed paragraph of the Hotel Reviews example

# DCNN as a decoder for reconstruction task

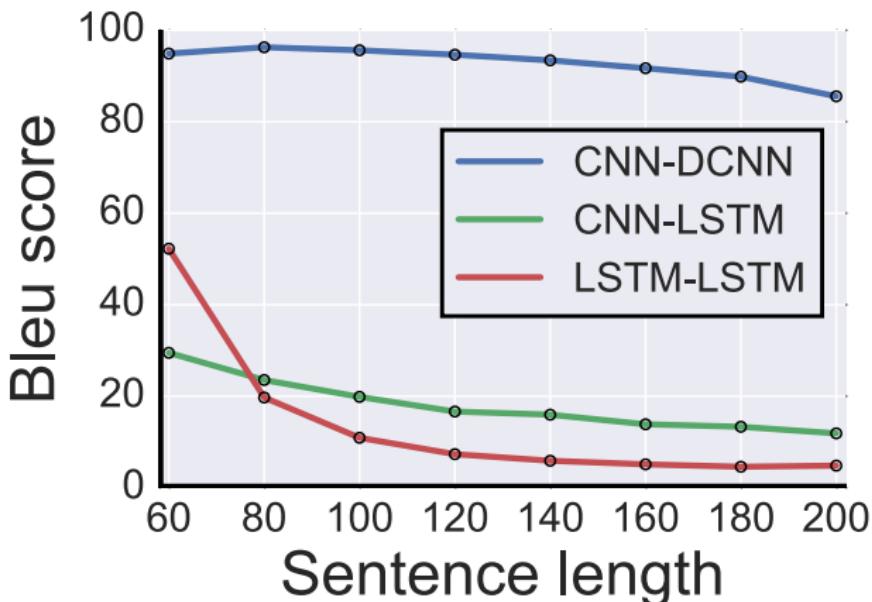
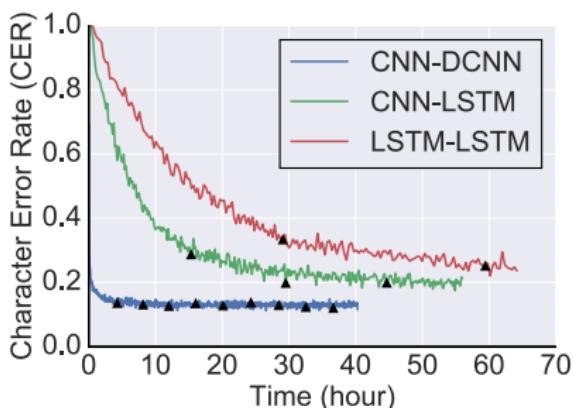


Figure: BLEU score vs. sentence length for Hotel Review data.

## Deconvolutional text autoencoder

## DCNN as a decoder for spelling correction task



Original	what s your idea of a stepping stone to better things to come ?
Modified	wuat s yogr idem of t stepukng jtzne ti better thingz tt coee ?
ActorCritic	what s your idem of t stepuang jokne ti better thing itt come ?
LSTM-LSTM	what s your idea of a speaking stand to better things to come ?
CNN-LSTM	what s your idem of a stepping start to better thing to come ?
CNN-DCNN	what s your idea of a stepping stone to better things to come ?

Figure: Spelling error denoising comparison.

## Deconvolutional text autoencoder

## DCNN as a decoder for title prediction task

Model	CNN-LSTM	CNN-LSTM w/ rec	CNN-DCNN	CNN-DCNN w/ rec
Rouge-L	16.37	<b>18.14</b>	14.75	16.83

Table: arXiv data. w/ rec: with DCNN reconstruction

Abstract:	this paper presents a new state of the art for document image classification and retrieval , using features learned by deep convolutional neural networks ( cnns ) . in object and scene analysis , deep neural nets are capable of learning a hierarchical chain of abstraction. experiments show that features extracted from cnns are robust to compression , cnns trained on non document images transfer well to document analysis tasks , and enforcing region specific feature learning is unnecessary given sufficient training data . this work also makes available a new labelled subset of the collection , useful for training new cnns for document analysis .
Ground-truth	evaluation of deep convolutional nets for document image classification and retrieval
CNN-LSTM	deep learning for image recognition
CNN-LSTM w/ r	a classification algorithm for deep neural networks
CNN-DCNN	deep neural convolutional for for recognition
CNN-DCNN w/ r	residual and based neural based classification for recognition net-

Deconvolutional text VAE

# Outline

## 1 Overview

## 2 Generating text via adversarial training

- Generating realistic text from given corpus
- Generating coherent and diverse conversation

## 3 Deconvolutional text modeling

- Deconvolutional text autoencoder
- Deconvolutional text VAE

## 4 Summary

## Deconvolutional text VAE

## Deconvolutional text VAE (on-going work)

- Leveraging variational auto-encoder to generate text continuously.
- For an input  $x$ , the distribution of latent code  $z$  is modelled as  $q_\phi(z|x)$ :

$$\begin{aligned}\mu &= g_1(f^{cnn}(x; \phi_1); \phi_1), \log \sigma = \mu = g_2(f^{cnn}(x; \phi_2); \phi_2) \\ z &\sim \mathcal{N}(\mu, \sigma)\end{aligned}$$

- The generative model  $p_\theta(x|z)$  is modelled as deconvolutional network.
- Variational lower bound

$$\mathcal{L} = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x)||p(z)) \quad (12)$$

## Deconvolutional text VAE

## Discussion

- Improves representation learning and downstream tasks

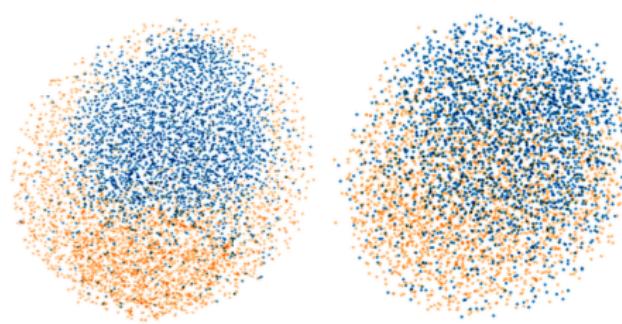


Figure: t-SNE embeddings (left: DCNN-VAE, right: LSTM-VAE) for BookCorpus and arXiv sentences, colored as orange and blue.

- Generation is less coherent comparing with LSTM.
- Have the best of both worlds to generate long sentences?
- *Globally* DCNN + *locally* RNN?

# Summary

- *A deconvolutional decoder*: performance is irrespective of length
- Significant computational savings.
- Better at representation learning.
- Still experimental for generation task.
- Future directions:
  - Could it be improved for language modeling/generation task by incorporating with RNN?
  - Following up, could it be used in GAN setup to avoid exposure bias issue?
- Code availability:  
[https://github.com/dreasysnail/textCNN\\_public](https://github.com/dreasysnail/textCNN_public)

# Conclusion

- Towards generating *consistent* and *diverse* text and alleviating issues of traditional RNN-based model, by leveraging DGMs.
- Text and conversion generation via *adversarial training*
- *Deconvolutional decoder* and its potential applications.

# Acknowledgment



Lawrence Carin  
(Duke)



Michel Galley  
(MSR)



Jianfeng Gao  
(MSR)



Bill Dolan  
(MSR)

**My collaborators:** Zhe Gan, Kai Fan, Dinghan Shen, Ricardo Henao, ...

# Thank You!